Autoreason — Experimental Results

Nous Research · https://github.com/NousResearch/autoreason

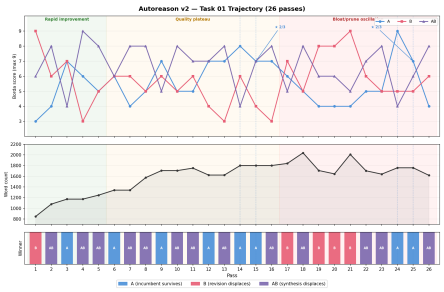

Figure 1: **Autoreason v2 trajectory (26 passes).** Top: Borda scores per pass. Middle: word count showing bloat/prune oscillation. Bottom: winner. Dashed lines mark 2-consecutive-A-win points.

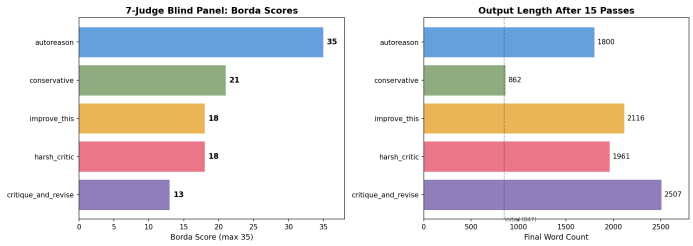

(a) **7-judge blind panel.** Borda scores (left) and word counts (right). Autoreason: 35/35, 7/7 first place. Conservative outranks all iterative baselines.

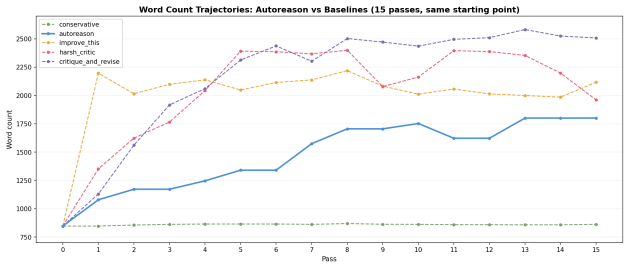

(b) **Word count trajectories.** Conservative flatlines, autoreason stabilizes, iterative baselines bloat. Critique-and-revise bloats most, scores worst.

Figure 2: **Baseline comparison.** Same start (847 words), same model, 15 passes. Judges see task, initial output, and all 5 finals (randomized labels).

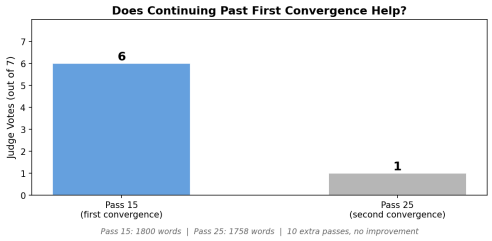

Figure 3: **Does continuing past first convergence help?** Pass 15 vs pass 25, 7-judge blind panel: 6–1 for pass 15. First convergence is the quality ceiling.